

The American Naturalist

The importance of measurement uncertainty in ecological management

--Manuscript Draft--

Manuscript Number:	58239
Full Title:	The importance of measurement uncertainty in ecological management
Short Title:	Measurement uncertainty
Article Type:	Major Article
Manuscript Classifications:	Conservation biology; Ecology: population; Fish; Modeling: population ecology; optimality theory; Population: dynamics
Additional Information:	
Question	Response
Please provide the word count for the main text (excluding the abstract, the literature cited, tables, or figure legends) [original submission]	8926

Response to reviewers

2018-01-31

ASSOCIATE EDITOR'S RECOMMENDATION

Dear Professor Bronstein; Prof Dwyer,

My recommendation is that manuscript 57722 be declined without prejudice. The authors use a state of the art computer algorithm known as "POMDP" to determine, on the fly, how fisheries should be managed, in terms of the total amount of stock that can be harvested, and the profit accruing from that harvest. This is an important problem in ecology, and the authors approach appears to offer significant insight. The reviewers nevertheless point out some important ways in which the text could be improved, and reviewer 1 in particular argues that the authors's approach may be significantly flawed. These problems would therefore need to be addressed for a revision to be successful.

We appreciate the careful and constructive comments of the editor and reviewers, which have allowed us to re-write the manuscript to address each of the issues raised here.

Reviewer 1's primary objection essentially re-asserts the "paradox of measurement uncertainty" that previous literature has also asserted: that uncertainty should somehow make less cautious. We appreciate the reviewer highlighting the importance of explaining this central issue properly, and our revisions should make it clear how this erroneous and potentially dangerous conclusion arises from simplifying assumptions. Reviewer 2's primary objection is the missed opportunity of focusing too narrowly on fisheries. We agree entirely with the reviewer that the fundamental issues here are germane to ecological management and not specific to fisheries, and have revised the presentation substantially to address this. Both reviewers raised concerns about the clarity of the methods; so in addition to rewritten discussion of methods we now provide fully executable code to replicate the analysis shown here. Lastly but most importantly, both reviewers raise the observation that few fisheries are actually managed using a strict constant escapement (CE) policy. We have introduced more realistic management strategies based on regulating fishing mortality through maximum sustainable yield (MSY) and total allowable catch (TAC), as additional comparisons. TAC illustrates a different, more heuristic approach to uncertainty than seen in the optimal control methods that provides some valuable additional insights into the challenge of accounting for measurement error. We think this results in a much more interesting manuscript which should be relevant to the broad audience of the American Naturalist.

Below we provide more detailed responses to each of the concerns raised and how they have been addressed in this revision.

One of Reviewer 2's main comments is that the article focusses too narrowly on fisheries and stock assessment, and the Reviewer argues that the authors could instead frame their work more broadly in terms of applied ecology in general, perhaps by citing problems of measurement uncertainty in other applications besides fisheries. These points are very well taken, especially given that The American Naturalist covers biological problems in a very broad sense. For the work to be of general interest to the readership of The American Naturalist, the authors would therefore have to explain more clearly the relevance of their work to applied ecology in general.

We now emphasize that making decisions in face of imperfect information is by no means an issue isolated solely to fisheries, which merely serves as a concrete example where relevant data can be easily obtained. We also note that fisheries management has long been a crucible in which ecological management theory and practice has been forged before becoming widely applied in other ecological contexts.

Reviewer 1 in contrast makes a specific and detailed comment about a very particular issue. In discussing this comment, an important point is that I am sufficiently far from being a fisheries biologist that I am not entirely

understand sure that I understand the objection, emphasizing the importance of Reviewer 2's comment that The American Naturalist is a general interest journal, not a fisheries journal. This caveat aside, my understanding of Reviewer 1's comment is that the Reviewer ran a set of simulations which were essentially the same as in the manuscript, except that the Reviewer explicitly excluded the possibility of 90 percent harvest of the stock population.

The analysis of Reviewer 1 is very interesting, and finds essentially the same paradoxical result as earlier published studies (e.g. Clark and Kirkwood (1986); Sethi et al. (2005)); as the reviewer summarizes:

In every case the optimal escapement goal was lower when observation error was included.

This is a paradoxical conclusion that defies the precautionary principle, and as such would have important ramifications in ecological management in many applications beyond fisheries. Further, though the reviewer's analysis is not identical to either the approaches of Clark and Kirkwood (1986) or Sethi et al. (2005) (neither assume constant escapement and both find this paradoxical result in only some but not all cases), we believe our analysis provides both a refutation of this result and a clear intuition as to where this misleading conclusion has come from.

As the editor notes, we suspect that Reviewer 1's simulations are not indeed identical except in the way mentioned, since they appear to assume the optimal policy can be specified in terms of a single constant escapement. Our results and in general escapement need not be constant but will depend on the observed stock size and on previous measurements, as we now emphasize. Also as the editor notes, the Reviewer's description of what they have done is extremely brief and thus difficult to be sure we have characterized precisely, but nevertheless, we believe our analysis sheds some light on why the reviewer finds this paradoxical conclusion while our analysis does not.

The escapement recommended by our POMDP solution is not constant, but depends both on prior information and the current estimate of the stock size, as illustrated in Figure 4. Assumptions that simplify the calculation of the policy by ignoring these aspects can thus reach different conclusions. Clark and Kirkwood (1986) assumes that the stock size is known perfectly at the end of the prior period, Sethi et al. (2005) implicitly assumes that the prior belief is always uniform rather than reflecting previous measurements, which puts erroneous probability weight on the chance the stock is in fact very high.

Figure 4 offers a clear explanation for when it is optimal to increase harvest in response to increased uncertainty: the POMDP harvest only exceeds the CE harvest when the measured stock size is erroneously low. That is, when the prior belief would have suggested the observation should be much higher, and therefore the measurement is most likely an error under-estimating the true stock size. This logic is entirely intuitive and is borne out in the resulting POMDP policies.

In the reviewer's simulations, the optimal escapement level is lower in the presence of measurement error, whereas in fig. 5 of the manuscript, the optimal escapement level is instead higher in the presence of measurement error. Reviewer 1's description of their simulations is brief enough that I would not expect the authors to build their entire revision around an effort to respond to this comment. If the authors have indeed allowed for 90 percent harvest, however, it is important that they test whether their results are different if they place an upper limit on the fraction of the stock that they allow to be harvested.

The reviewer and editor have raised an important issue in the question of whether we permit the harvest quota to be set to a value higher than the stock. There are two problems with this. First, it's not clear how the reviewer implements this assumption in the face of measurement uncertainty: how can one be certain a quota is less than 90% of stock size X without knowing what X is? One can obviously enforce this with respect to the measured stock size but not the true stock size. More importantly, this assumption amounts to the assertion that it is mathematically impossible to over-fish a stock to the point of collapse. However low the stock size might get, local extinction is impossible. In this world-view, there will always be some tiny fraction of fish from which the population can recover. Assuming that local extinction is impossible will indeed bias any optimal decision model to be more aggressive. Obviously this also has implications for ecological models in general, conservation policy would look different if we assumed extinction was impossible. Clark and Kirkwood (1986) already recognized that a very similar assumption in Reed (1979) derivation

that stochastic models needed no additional caution ($S = D$) was likely due to the fact that accidentally extinction was impossible (see manuscript, line 150). We do not adopt this assumption. Obviously real world populations cannot be indefinitely reduced by only a fractional percentage. In our simulations, though we consider a Gordon-Schaefer model, stock abundances are divided into 100 distinct bins, with the smallest bin corresponding to a stock size of zero. Thus it is possible for over-harvesting or stochastic environment, or some combination of both, to lead a population to permanently collapse once it has gotten small enough.

I also agree with Reviewer 1's comment that the methods are not that easy to follow. That said, to my understanding, the escapements in fig. 2 are not constant, and so Reviewer 1's comment that a constant escapement cannot lead to the collapse of a stock seems to me to be irrelevant, but again I may be wrong, and my uncertainty emphasizes the lack of clarity in the writing, especially with respect to the Methods.

We agree entirely with this, and hope that the rewritten description, together with complete code, helps make this somewhat clearer.

Page 8-12 in particular are written in a very formal style that I associate with journals that focus on mathematics rather than on biology, with the difficulty problem that much of this text has small grammatical errors that exacerbate the clarity problem. Meanwhile, both Reviewers argued that the authors have not accurately characterized the literature in this area, and that problem would also have to be solved in a revised version.

The presentation of all of this material has been entirely re-written in what we hope you will find is a much more appropriate and accessible style. To avoid needless formalism, we do not rehash the technical details of solving Partially Observed Markov Decision Processes but instead point to the extensive literature already existing on the subject. Instead, we focus on describing the nature of the problem and how it differs from analyses such as the perfectly observed Markov Decision Process that provides the basis for the Constant Escapement strategy (Reed (1979)) and from other Bayesian hierarchical modeling (e.g. partially observed Markov processes, pomp, or Hidden Markov Models, HMMs) which are not decision processes at all, but rather a prerequisite part of choosing & parameterizing a model, not using said model to make a decision.

In spite of these negative comments, it seems like the overall approach has a great deal of potential, as Reviewer 2 argues, and if Reviewer 1's concerns prove to be minor, then I believe that a revised version would be a substantial contribution. That said, the complexity of the technical issues is quite high, and so it may be simpler to submit the manuscript to a journal that specializes in mathematical biology or in fisheries management. – Greg Dwyer

Reviewer #1:

For complete review, see the attached PDF, which begins as follows:

I attempted to evaluate the papers overall conclusions by a very simple simulation and these results seem to challenge the papers conclusions. I simulated a population with the model the authors describe with both observation and process error lognormally distributed and using the authors parameters. For a range of escapement goals I calculated the net present value of using that escapement goal, and compared the consequences when there was no observation error to the case where the lognormal cv was 0.5 – that is very high observation error. The one difference I made was I did not let the harvest exceed 90% of the total population size ... this seems reasonable as I know of few fisheries where 90% of the population could be harvested in a single year.

The graph below shows a typical result. The thin line is with no observation error, the thicker line with the

observation error. In every case the optimal escapement goal was lower when observation error was included.

We thank the Reviewer for this careful parallel investigation, which was very useful in focusing the manuscript on some of the key results. We believe the differences in our respective results can be explained by the differences in our methods as discussed in our reply to the Editor’s discussion of this issue. In short, we believe the reviewer’s seemingly counter-intuitive result that increased uncertainty should be met with an action (lowering escapement) that would increase the risk of stock collapse rather than the intuitive response to be cautious by reducing harvest, is simply an artifact of the reviewer’s assumption that stock collapse is impossible, and thus increased harvesting poses no real risk.

I can’t really understand some of the results in the paper ... for instance I don’t see how a fixed escapement policy could lead to the collapses shown in Figure 2 – using fixed escapement policies you can’t have you population continue to decline.

We agree that the results and methodology could have been much more clearly presented. We hope the rewritten presentation of these results clarifies these issues. In particular, note that once you introduce measurement error a fixed escapement policy can only achieve a fixed *expected* escapement, you never know what the actual escapement is because you never know exactly how many fish you started with. As measurement errors are equally likely to be too high or too low, a run of “heads”, so to speak, would result in consistently over-harvesting and continued decline (see discussion around line 363.)

Finally it is certainly not true that fixed escapement policies are commonly used. To my knowledge almost all marine fish populations are managed with policies more resembling fixed harvest rates with caveats that target harvest rate declines when stocks drop below a target stock size. Certainly some salmon stocks are managed with an escapement goal, but they are not in any sense common or the default in marine fisheries management.

The reviewer is entirely correct about this, and we greatly appreciate the reviewer raising this important issue. In addition to the previous analysis which focused only on comparison to a constant escapement policy, we now also include comparisons to a more realistic “Total Allowable Catch”, or TAC strategy.

Reviewer #2:

In this clearly written paper the authors show that fish population dynamic models that ignore measurement error can result in significant overharvesting, decreased economic returns, and the potential for complete collapse of the stock. The key advance is the use of an algorithm for partially observed Markov decision processes (POMDP) to provide a complete solution rather than relying upon inferences from a deterministic model or a model of optimal control that does not update belief states during each time step.

The treatment of this issue is remarkably comprehensive with respect to structural uncertainty of the model formulation. In fact, the authors in some way bury what I see as the lead in the Appendix. They include analysis of a population subject to an Allee effect, and produce the same general conclusion as they do via analysis of models with more standard dynamics: the POMDP approach is less likely to cause overharvest and population collapse than alternative approaches, even if the location of the unstable steady state population size (tipping point) is unknown. This finding is very important given the increasingly voluminous literature on thresholds and tipping points, suggesting that proper accounting for measurement uncertainty may be a practical alternative to investing more in reducing such uncertainty in the budget-constrained world of natural resource management.

As is clear from my summary, this paper is framed for a narrower audience focused on fisheries and stock assessment. To draw in a broader readership, I suggest that the authors motivate the problem they address and interpret their results in the context of the more general ecological, wildlife, and forestry literatures. Specific examples of how measurement uncertainty has been found problematic in fisheries and wildlife assessments would be helpful.

We thank the reviewer for this careful and constructive review, and agree entirely with the recommendation that the fundamental issues here are not limited to the problem of fishing quotas but are germane to much of ecological decision making. We have now tried both to set up the generic nature of the issue of measurement uncertainty (see Introduction) and discuss its implications beyond fisheries much more explicitly (e.g. lines 423 - 437). We have however removed the treatment of the tipping point model in order to focus explicitly on the role of uncertainty in the more familiar case of logistic growth, where the dramatic consequences of measurement uncertainty may be less intuitive and cannot be dismissed merely by questioning the frequency of tipping points occurring in real systems. We hope to address the problem of tipping points more explicitly and more fully in later work.

The authors mischaracterize the extent to which measurement uncertainty is left untreated in stock assessments. Dichmont et al. 2016 (Fisheries Research) provide a nice review of how some of these measurement uncertainty issues are treated in practice in US stock assessments, and an entry point to the literature. If the authors intend to focus on theoretical models and general solutions for population dynamics, their assertion about failure to incorporate measurement uncertainty completely may be valid but this should be clarified.

We had no intention of suggesting that measurement uncertainty was under-represented in *stock assessments*; indeed it is precisely the *prevalance* with which stock assessment models account for uncertainty that motivates this work. We have extensively revised the manuscript to be much clearer on this point. Our assertion is that these estimates of the uncertainty associated with a measurement of the stock size to which the reviewer refers are not properly accounted for in the process of setting a quota.

Unlike the classic pretty good yield paper mentioned in the paper (line 440), the authors do not adequately address how conservative harvest is likely to be met with opposition by harvesters and how this issue can be addressed directly by their analysis.

This is an excellent point which we now develop in the discussion, lines 456-470.

Similarly, rather than dismiss management strategy evaluation as a less preferable alternative to the optimal control approach taken here, it would be nice to see the authors include a formal contrast. Might the offline POMDP approach be framed this way and results inferred without additional analyses? This would be nice because in the real world, optimal control with annual updating of harvest regulations (the POMDP approach) may not be practical. The offline POMDP seems more realistic, perhaps with an online phase implemented such that belief states could be updated in 5-10 year intervals.

We now include a formal comparison against TAC-based management policy (and also MSY based policy in the supplement), which reflect this aspect of today's actual management policies which are not updated annually but apply a constant fishing effort (mortality). We also discuss several of the practical objections to implementing POMDP in practice and how it might be modified to address more of these concerns, e.g. lines 456-490.

Minor comments: The analysis related to Fig 5 made me wonder how POMDP compares to alternative frameworks when a stock is increasing or decreasing in abundance. The authors mention this in their interpretation, but I think it could be highlighted more as it seems that the POMDP approach (at least for plaice) allows a more careful tracking of catch to stock size during collapse and rebuilding, a result that would be viewed favorably by fishermen and conservationists alike.

Excellent observation! Yes, due to the influence of the prior belief, POMDP is slightly more cautious when a stock is rebuilding from small sizes. We now speak to this in our discussion of Figure 5. We agree with the reviewer’s general observation that the POMDP approach allows for more careful tracking of stock size despite measurement error, which in turn leads to both faster rebuilding of stocks and higher economic yields.

Some reorganization may be warranted. Several portions of the Results section describe Methods, which would be nice to signal earlier in the MS (e.g., lines 266-281, lines 298-300 and lines 321-324 could be included in Methods to clarify that the POMDP overestimates measurement uncertainty, etc. lines 348-351, etc).

We agree entirely, and now present the full discussion of the methods and models before going into the results.

In several places the authors use loose language by referring to marine ecosystem dynamics when in fact they are treating population dynamics.

Agreed, this has been addressed in the rewrite.

The authors discuss volatility of different solutions and would be served well to quantify it. E.g., lines 365-369

We believe this is much better illustrated by the visualization of these solutions in Figure 5, and introducing a statistic to quantify this would only needlessly confuse that discussion.

References

- Clark, Colin W, and Geoffrey P. Kirkwood. 1986. “On uncertain renewable resource stocks: Optimal harvest policies and the value of stock surveys.” *Journal of Environmental Economics and Management* 13 (3): 235–44. doi:10.1016/0095-0696(86)90024-0.
- Reed, William J. 1979. “Optimal escapement levels in stochastic and deterministic harvesting models.” *Journal of Environmental Economics and Management* 6 (4). Elsevier: 350–63. doi:10.1016/0095-0696(79)90014-7.
- Sethi, Gautam, Christopher Costello, Anthony Fisher, Michael Hanemann, and Larry Karp. 2005. “Fishery management under multiple uncertainty.” *Journal of Environmental Economics and Management* 50 (2): 300–318. doi:10.1016/j.jeem.2004.11.005.

The importance of measurement uncertainty in ecological management

Authors removed during peer review

3 Abstract

Ecological management and decision-making typically focus on uncertainty about the future, but surprisingly little is known about how to account for uncertainty of the present: that is, the realities of having only partial or imperfect measurements. Our two primary paradigms for handling decisions under uncertainty – the precautionary principle and optimal control – have so far given entirely contradictory results. This paradox is best illustrated in the example of fisheries management, where many ideas that guide thinking about ecological decision making were first developed. We find that simplistic optimal control approaches have repeatedly concluded that a manager should increase catch quotas when faced with greater uncertainty about the fish biomass. Current best practices take a more precautionary approach, decreasing catch quotas by a fixed amount to account for uncertainty. Using comparisons to both simulated and historical catch data, we find that neither approach is sufficient to avoid stock collapses under moderate observational uncertainty. Using partially observed Markov decision process (POMDP) methods, we demonstrate how this paradox arises from flaws in the standard theory, which contributes to over-exploitation of fisheries and increased probability of economic and ecological collapse. In contrast, we find POMDP-based management avoids such over-exploitation while also generating higher economic value. These results have significant implications for how we handle uncertainty in both fisheries and ecological management more generally.

Key words: Measurement uncertainty, POMDP, Decision theory, Fisheries, Conservation,

Imperfect information is ubiquitous in ecological management and conservation decision making. While the pressing concerns of global change have put the spotlight on *forecasting*, i.e. the uncertainty of the future (e.g. Petchey et al. 2015), management decisions must also contend with uncertainty of the *present*: How many fish are in the sea today? What regions harbor the most fragile biodiversity? How have invasive species transformed native communities? A recognition of the importance of uncertainty as well as sophisticated methods and tools to do so has become increasingly widespread in the context of estimating ecological models from available data, thanks in large part to the rise of approaches such as Hierarchical Bayesian Modeling, along with convenient software tools for implementing these methods. In sharp contrast to that shift, the treatment of uncertainty in decisions – how the uncertainty in models should be reflected in scientific

recommendations, remains far less developed, more qualitative, and under-served by available tools. In this
15 paper we illustrate how uncertainty in measurement has a dramatic negative impact on outcomes under
a variety of these existing approaches, which can be avoided through the use of more formal approaches
that have since been developed in other fields that can likewise be implemented in software tools. This
18 insight both resolves a long-standing paradox in fisheries management and has important implications for
how ecosystem management deals with uncertainty more generally.

Approaches for decision making under uncertainty in ecological systems can be divided into two camps:
21 approaches based in “optimal control” and usually favored by natural resource economists, and approaches
based in heuristic methods such as scenario planning, resilience thinking, and precautionary rules of thumb,
more commonly found in both ecological literature and actual practice (Fischer et al. 2009; Polasky et al.
24 2011). While researchers have for some time recognized the need to unify the transparent and quantitative
algorithmic approach of optimal control with the greater complexity and uncertainty of real ecosystems
that is acknowledged by heuristic methods, computational barriers to doing so have hither-to stymied this
27 progress. Here, we illustrate how the limitations of these approaches has been manifested in the example of
fisheries management, and present a new approach that can combine the reality of measurement uncertainty
and the rigor of optimization to resolve a long-standing paradox and suggest a more robust approach to
30 management.

Fisheries conservation and management has long been both a crucible and proving ground for the theory
of ecological management more generally, including topics such as adaptive management (Walters and
33 Hilborn 1978), ecosystem-based management (Levin and Lubchenco 2008) and resilience thinking (Holling
1973; May 1977), while also giving rise to the sub-discipline of resource economics (Gordon and Press 1954;
Schaefer 1954; Beverton and Holt 1957), thanks to its global relevance, long history, and readily available
36 data. Consequently, understanding the impact of measurement uncertainty in this rich and well-studied
context will have implications for both resource management and conservation efforts in other domains, many
of which draw on concepts established or tested in fisheries ecosystems. In particular, the approach here
39 illustrates the limitations of both over-simple optimal control approaches and rule-of-thumb precautionary
approaches, and offers instead a more general strategy for dealing with measurement uncertainty adaptively.

We compare the results of managing simulated fish stocks under the existing optimal control approach,
42 a heuristic approach that more closely resembles today’s management practices, and a newly proposed
approach that provides a more rigorous treatment of uncertainty.

The Decision Problem

We consider the management problem of setting catch quotas for a marine fishery in the face of imperfect information about the current stock size and uncertainty about future recruitment.

We seek to determine the sequence of actions h_t for $t \in [1, \dots, \infty]$ that maximize the net present value (discounted sum of all future profits) of the fishery. We will denote the discount rate γ . For simplicity we will once again follow classic theory and assume a fixed price for fish, (equivalently, measuring our value in units of discounted fish rather than discounted dollars), though again this assumption is easy to relax. Each year, the manager also obtains an estimate y_t of the true x_t stock size subject to some measurement uncertainty ξ_t

$$y_t = \xi_t x_t$$

For simplicity we will assume measurement uncertainty is normally distributed with standard deviation σ_m , $\xi \sim \mathcal{N}(1, \sigma_m)$. Given an estimated recruitment model along with this measurement of the stock size, the manager must choose the set of a harvest quotas $\{h_t\}$ for $t \in [0, 1, \infty]$ that maximizes expected long term (net) present value:

$$\max_{\{h_t\}} \mathbb{E} \left\{ \sum_{t=0}^{\infty} \gamma^t \cdot h_t \right\}$$

A few subtleties arise in this seemingly simple problem statement that we must address. First, given our discrete-time formulation of the recruitment process, it is necessary to decide if the measurement y_t happens before or after harvest h_t : that is do we: measure, recruit, harvest, or measure, harvest, recruit? Following convention in the optimal control literature (e.g. (Reed 1979; Clark 1990; Sethi et al. 2005)), we will assume the latter; measurement occurs immediately before harvest, followed by recruitment.

A second subtlety is in the nature of the decision problem itself: note that the manager is not myopic, thinking only about the next year, but rather considers the sequence of all future actions h_t . While the manager will revisit this sequence each year in light of new observations, it is this ability of the decision problem to think ahead that allows it to tolerate short-term costs (e.g. reduced harvests in the current season) for larger future payoffs. In this, ecological management is a game of chess, always thinking several moves ahead. This differs from the problem solved in theory as well as the practice of maximum sustainable yield, which seeks to identify only a fixed mortality F_{MSY} applied for all time, rather than a *policy* of varying quotas h_t depending on the stock assessment y_t . Understanding this distinction is key both to understanding the differences between current theory and current best practice in fisheries management.

Population model

To facilitate tractability and interpretation, we will focus on the well-studied Gordon-Schaefer model (Gordon and Press 1954; Schaefer 1954) of logistic population growth (stock recruitment),

$$x_{t+1} = \varepsilon_t(x_t - h_t)r \left(1 - \frac{(x_t - h_t)}{K}\right)$$

where x_t is the current stock size, h_t the harvest chosen that year, with parameters r giving the individual growth rate, and K the carrying capacity and ε representing stochastic recruitment in a variable environment.

We have written this in terms of $x_t - h_t$ to underscore that we assume the convention of observe, harvest, recruit. We will assume for simplicity $\varepsilon \sim \mathcal{N}(1, \sigma_g)$. This and similar models are widely used in large scale analyses across diverse stocks (e.g. (Costello et al. 2016; Britten et al. 2017)), and form the basis for much of bioeconomic theory (Clark 1990). As such, it will be easier to compare our results against intuition and classic theory, and avoid the possibility of differences arising only because of some particular subtle assumption hidden in a more complex model.

It is important to remember that both existing methods we discuss and the POMDP approach proposed here can or are already applied to more complex models as well, including those with age or stage structure. Likewise, while we consider the model and it's parameters as "given," it should be understood to mean that this in practice comes as the result of some model choice and statistical estimation from historical data, frequently accounting for uncertainty in both the process and the measurements (e.g. Dichmont et al. 2016). As we shall see here, having included uncertainty in the model *estimation* in no way excuses us from having to also deal once again with that uncertainty in the decision process.

Current Theory and Practice

Gordon and Press (1954) and Schaefer (1954) independently showed in the same year that the maximum sustainable yield for the model bearing their names is achieved by reducing the stock to the population to the size at which it obtains the maximum growth rate which is known as Biomass at Maximum Sustainable Yield, B_{MSY} . For the Gordon-Schaefer model (and many others), this is achieved at $B_{MSY} = K/2$. They observed that fishing at a *constant yield* (that is, an individual fish mortality, F , such that harvest $H = F \cdot B$) of $F_{MSY} := \frac{H_{MSY}}{B_{MSY}}$ will eventually lead to a population that converges to the biomass B_{MSY} and produces the maximum sustainable harvest, $H_{MSY} = rK/4$ in this model. From here, theory and practice diverge.

This *constant yield* (constant mortality) solution thus corresponds to an equilibrium analysis which does not solve the time-dependent optimization problem above. In particular, this maximum sustainable yield

(MSY) strategy will not be optimal whenever the stock is away from B_{MSY} . Clark (1973) demonstrated (assuming no uncertainty in measurement or stochasticity in population dynamics) that the optimal time-dependent strategy is not one of constant yield, but rather of *constant escapement*. We will return to this approach in a moment.

Uncertainty in Practice: MSY & TAC

Maximum Sustainable Yield (MSY) remains the basis of international law (including the UN, IWC, IATTC, ICCAT, ICNAF; Mace (2001)) and a familiar standard of management in terrestrial ecosystems as well as aquatic (Clark 1990). Critics have for some time observed the limitations of harvesting at MSY in face of uncertainty in stock sizes and population dynamics (Larkin 1977; Botsford et al. 1997). Many US fisheries reflect this uncertainty through a series of adjustments that effectively reduce the target fishing mortality level to reflect this uncertainty. Typically, a stock assessment model provides a best estimate of B_{MSY} and corresponding mortality F_{MSY} is used to define the stock Over-fishing Limit, (OFL). Based on this, a somewhat lower level is set as the Allowable Biological Catch (ABC), reflecting uncertainty in the stock assessment. To reflect possible uncertainty between reported and actual catch, the ABC is reduced somewhat further to define the Total Allowable Catch, (TAC), which forms the basic unit of management for many such fisheries. To reflect this process, our analysis will also consider policies in which the harvest quota is set at 80% of the level expected under the MSY policy: $H_{TAC}(t) = 0.8 \cdot F_{MSY} \cdot B_t$, for a biomass estimated at $B(t)$ in year t . To distinguish this approach from MSY, we will refer to this approach as a TAC policy. This more closely represents a heuristic (e.g. Hilborn 2010) or resilience-based approach (sensu Fischer et al. 2009) than an optimization-based policy. Importantly, this approach shares the fundamentally stationary assumptions of an MSY policy by defining a constant mortality rather than a dynamic policy. Thus, despite being more cautious overall and generating lower economic yield than a constant escapement policy, TAC policies continue to harvest at non-zero rates even if a stock falls below B_{MSY} , while the constant escapement policy does not.

Optimal management

So far, attempts to provide a more formal basis for managing uncertainty than the heuristic adjustment of catch limits described above have largely foundered in a series of paradoxes. Without uncertainty, the theoretical policies are quite intuitive: The result derived by Clark (1973) for an optimal dynamic strategy to replace the equilibrium solution of MSY can be summarized as: If the stock is most productive at B_{MSY} , then obtain B_{MSY} as quickly as possible. Thus, at any level below B_{MSY} , the economically optimal thing to do

is to completely shut down the fishery, while above this stock, harvests greater than H_{MSY} will be needed to
129 bring the stock back to B_{MSY} . (This intuition must be adjusted slightly in the case of economic discounting
of future profits, but is otherwise quite general, see Clark (1990)). This strategy is known as constant
escapement, since a constant stock size B_{MSY} escapes harvest each year. While both converge to the same
132 long-term biomass and long-term yield, these differences away from equilibrium between the mortality-based
approach that dominates fisheries management practice and the escapement-based approach that became the
focus of theory will result in different behaviors under uncertainty. Intuitively, any uncertainty, either in the
135 population growth model (stochasticity), or measurement error, should seem to justify harvesting fewer fish,
as is built into the TAC limits discussed above. Yet formal treatments of these issues have so far required
approximations, and these approximations have driven paradoxical conclusions which we briefly summarize
138 here.

Reed's Paradox: $S = D$

The first of these we shall refer to as Reed's Paradox for stochastic growth, owing to the mathematical
141 proof provided in Reed (1979) which demonstrates that, under sufficiently general assumptions, the optimal
escapement S for a population under stochastic growth, is identical to the optimal escapement D of a
deterministic population, $S = D = B_{MSY}$. This surprising result suggests that in going from a world where a
144 manager has perfect knowledge of absolutely everything into a scenario where the manager faces considerable
uncertainty about the future state of the world, no additional precaution is needed. This also provides the
manager with a remarkably convenient mechanism for determining the optimal harvest policy: rather than
147 rely on computationally intensive Stochastic Dynamic Programming to determine the optimal policy, a
simple derivative will suffice, which is all that is needed to solve the maximum of the deterministic problem.

Clark's Paradox

150 Clark and Kirkwood (1986) was among the first attempts to resolve Reed's Paradox. Clark and Kirkwood
(1986) (quite correctly, as we will see), identified the crux of Reed's Paradox as the absence of measurement
uncertainty:

153 An important tacit assumption in Reed's analysis, as in the other works referred to above, is that
the recruitment level X is known accurately prior to the harvest decision, [...] In the case of
fishery resources, the stock level X is almost never known very accurately, owing to the difficulty
156 of observing fish in their natural environment.

Unfortunately, Clark and Kirkwood (1986) was unable to solve the resulting problem exactly, but had to adopt an almost equally troublesome assumption:

For reasons of tractability, we shall adopt the simplifying assumption that the escapement level S , is known exactly at the end of that period. (The mathematical difficulty of the problem increases markedly if this assumption is relaxed.)

Unfortunately for Clark and Kirkwood (1986), this “simplifying assumption” serves to squash most of the measurement error, and their results instead only deepened the paradox, finding policies that become even *less* cautious as uncertainty increases:

[Our] results appear to contradict the conventional wisdom of renewable resource management, under which high uncertainty would call for increased caution in the setting of quotas.

Relying on a quite different but still flawed assumption nearly two decades later, Sethi et al. (2005) largely confirm Clark’s Paradox, which they likewise observed with some concern:

It may seem counter-intuitive that a measurement error causes lower expected escapements below the deterministic fishery closure threshold.

Despite these notes of caution, both Clark and Kirkwood (1986) and Sethi et al. (2005) ultimately attempt to rationalize this counter-intuitive conclusion rather than reject it. Others have also attempted to introduce measurement error, but always relying on various approximations that similarly alter the problem. The most common such assumption is to assume from the outset that the solution must be of ‘constant-escapement’ type (Ludwig and Walters 1981; Roughgarden and Smith 1996; Engen et al. 1997; Moxnes 2003). Here we will illustrate the use of powerful modern algorithms that will permit a more direct solution to the case of measurement uncertainty. In addition to resolving this paradox, we will demonstrate this direct solution also out-performs current more heuristic rules of thumb such as those currently used in fisheries management to account for uncertainty.

POMDPs: An optimal treatment of measurement uncertainty

As discussed above, absent the short-cut provided by Reed’s theorems – solving for the optimal management policy in a stochastically fluctuating population – involves a Markov Decision Process (MDP), commonly referred to in ecological literature by its solution method, Stochastic Dynamic Programming (SDP). The constant escapement (CE) strategy considered here is the result of such an MDP or SDP analysis. Once

measurement uncertainty is introduced, the problem becomes a Partially Observed Markov Decision Process (POMDP, not be confused with a partially observed or hidden Markov model, HMM, in which no decision process is involved). The POMDP problem is posed identically to that of the MDP problem, with the addition of an observation process, but cannot be solved using SDP directly. The POMDP problem for fisheries question considered here can be summarized as follows:

- Transition function (state equation): $T(x_t, x_{t+1}, a_t)$: the probability that a system is in state x_{t+1} at time $t + 1$ given that it began in state x_t at time t and the manager took action a_t . In our context, this relationship is given by the Gordon-Schaefer stock recruitment function f with normally distributed growth uncertainty $x_{t+1} \sim \mathcal{N}(f(x_t, a_t), \sigma_g)$, truncated at zero to exclude negative population sizes.
- Observation function: $O(x_t, y_t, a_{t-1})$ the probability of observing state y_t given a system in state x_t . In principle, the action chosen can influence the precision of the observation. In our case, we simply assume normally distributed errors around the true state, $y_t \sim \mathcal{N}(x_t, \sigma_m)$, truncated at zero to exclude negative population sizes.
- Utility function: $U(x_t, a_t)$, the utility received at time t for taking action a_t , given that the system is in state x_t . For simplicity of analysis, we will simply set the utility to be equal to the harvested stock: $U(x_t, a_t) = \min(x_t, a_t)$, indicating that realized harvest cannot be negative. This choice ensures that in the case of no uncertainty ($\sigma_g = \sigma_m = 0$), the optimal solution matches that expected under a simple MSY calculation. More realistic utility functions may include diminishing returns with increasing harvest (supply and demand effects), and the cost of fishing, both of which act to suppress large harvests. By focusing on a simple utility we can be sure that our comparison to MSY is driven by the treatment of uncertainty rather than merely differing economic assumptions.

The optimization problem is to select the action a_t that will maximize the net present utility over all time. Future utility may be discounted by a factor γ , so that a value V in t year is valued at $\gamma^t V$ today. Numerically, each of these functions are defined over a discrete set of possible states, observations, and actions, and can thus be represented as a collection of matrices or tensors. While the solution method for MDP problems, stochastic dynamic programming (SDP) is well known and readily implemented in ecological problems (e.g Mangel and Clark 1988; Marescot et al. 2013), the POMDP problem cannot be efficiently solved using SDP. While not unknown to the conservation literature (Williams 2011), algorithms for POMDP have historically scaled quite poorly, and their application has been restricted to contexts with only a handful of possible states and actions (e.g. Chadès et al. 2008, 2011; Fackler and Haight 2014; Fackler and Pacifici

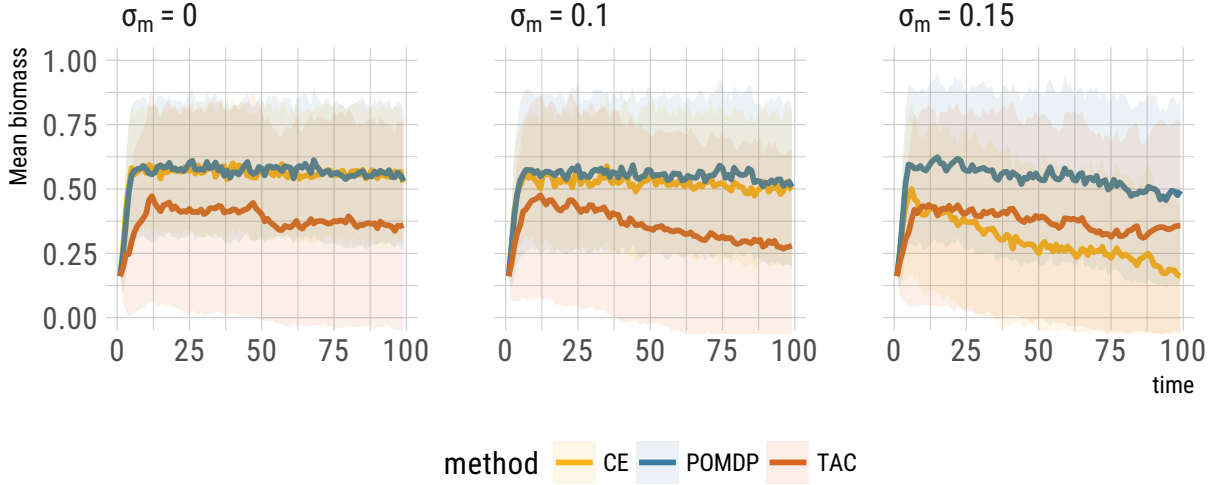


Figure 1: Average fish biomass under different management strategies under increasing levels of measurement uncertainty. Each plot the mean stock size over time across 100 replicate simulations under each policy: constant escapement (CE), Total allowable catch (TAC = 80% MSY), and the proposed partially observed Markov decision process (POMDP) method. Measurement error increases as a normal distribution with standard deviation 0, 0.1, or 0.15, as indicated at the top of the panel. Environmental stochasticity is fixed a standard deviation of 0.15 in each panel. Carrying capacity K normalized to 1, $r = 0.75$. Additional environmental noise levels and comparison to MSY rather than TAC can be found in the supplementary material.

2014). POMDP algorithms have a substantial literature starting almost contemporaneously to the work of Reed (1979) (see foundational papers by Smallwood and Sondik (1973) and Sondik (1978)) and remain an active area in artificial intelligence (e.g. Kaelbling et al. 1998; Pineau et al. 2003). By adapting recent algorithmic developments (Kurniawati et al. 2008) from that field we are able to find solutions for the considerably more complex problems such as the case of fisheries management, which requires on the order of 100 states and actions to provide sufficient numerical resolution. As the details for solving such problems numerically are already well documented we will not rehash them here, but instead provide a preliminary R package (Boettiger et al. 2018) which provides an implementation of the SARSOP algorithm and the fisheries POMDP problem, which can be used to replicate any of the results presented here and explore other variations. Details of the analysis presented here including annotated code to reproduce and further explore all of the following results is provided in the supplementary material.

Results

Figure 1 shows average fish biomass across 100 replicate simulations under three different management strategies: constant escapement (CE), total allowable catch (TAC; equal to 80% MSY), and partially observed Markov decision process (POMDP) management. Each successive panel shows a subsequently higher level of measurement error, from $\sigma_m = 0$, $\sigma_m = 0.1$, and $\sigma_m = 0.15$, as indicated. Standard deviation from the mean across replicate simulations is shown as faint colored bands, indicating significant variation due

to stochasticity between individual replicates. In each panel shown in Figure 1 stochastic recruitment (environmental noise) is set to a moderate $\sigma_g = 0.15$. In the absence of either environmental noise or measurement error (Supplemental material, Figure S1) POMDP, CE, and MSY would converge to stock at the B_{MSY} , while TAC would maintain the stock at a slightly higher level. The first panel of Figure 1 shows that the introduction of stochastic growth has a significant negative impact on the TAC strategy (This impact is even more severe for MSY, which always harvests strictly more than TAC by definition, as, shown in Supplemental material Figure S2), which results in an average biomass significantly lower than B_{MSY} . Without measurement error, the CE and POMDP strategies are nearly identical with both approximately maintaining the stock at B_{MSY} despite the significant environmental stochasticity. As measurement error increases in the subsequent panels, CE and TAC strategies perform increasingly poorly, while the POMDP continues to maintain the average stock close to B_{MSY} . Notably, CE is significantly more impacted by measurement uncertainty than TAC, with CE averaging even lower biomass than TAC under moderate measurement uncertainty.

Figure 1 confirms that the precautionary approach represented by the TAC does indeed prove more robust to the problem of measurement error than the optimization solution represented by CE. In general, we expect that the rigid assumptions required of optimization will lead to worse outcomes when those assumptions are not met than we see under the corresponding heuristic approach. Yet it is important to bear in mind that the opposite pattern was observed in the case of environmental stochasticity, in which it was the heuristic TAC strategy rather than the CE strategy that lead to the highest over-fishing rate. In contrast, the POMDP approach successfully handles both stochasticity and observation error, maintaining the stock near B_{MSY} levels that produce the highest yield.

This pattern in management success in ecological terms (the relatively recovery and maintenance of fish biomass) is also borne out in terms of economic performance. Figure 2 shows the mean net present value (averaging across replicates and discounting future profits by the discount factor γ) for these same simulations at increasing levels of measurement uncertainty. As before, environmental stochasticity is set at $\sigma_g = 0.15$. In the absence of measurement uncertainty, constant escapement (CE) is optimal (as per Reed (1979)), while the heuristic caution built into the total allowable catch (TAC) strategy (at 80% MSY) results in a sub-optimal economic yield. The presence of stochasticity in recruitment ($\sigma_g = 0.15$), also contributes to the reduction in yield under TAC. Though CE is quite robust to the environmental stochasticity level, this strategy proves very sensitive to increasing measurement error, showing sharp declines in economic value, falling below the TAC economic value at $\sigma_m = 0.15$. Though the risk of stock collapse does increase

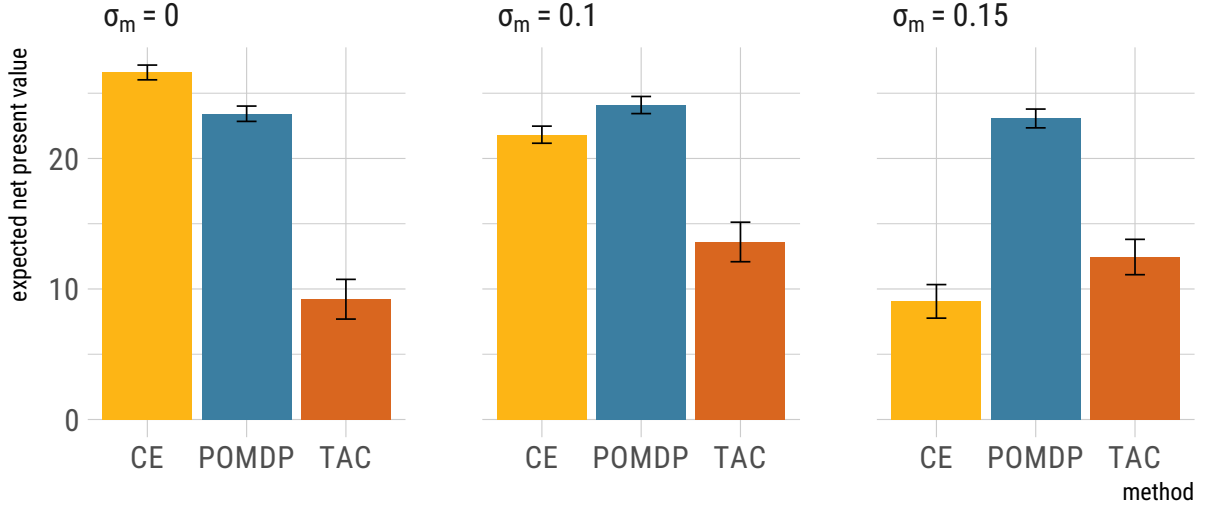


Figure 2: Expected net present value of the fishery across strategies under increasing levels of measurement error σ_m . Net present value is the sum of all future harvests discounted over time and averaged across 100 replicates simulations under each strategy, as shown in Figure 1. The value under constant escapement (CE) is optimal when measurements are perfect but decreases rapidly with increasing measurement error. The more cautious TAC is not economically optimal but largely unimpacted by increasing error, while POMDP attains consistently high economic yield despite the increasing uncertainty.

with increasing measurement error under TAC (as seen by the mean declines in Figure 1), these have little impact on the economic value due to the discount rate. A smaller discount rate would penalize unlikely but not improbable collapses more, since a significant amount of time is required to realize those rare events. Supplemental Figure S3 summarizes these economic trends across different noise values σ_g and includes comparison to a simple MSY policy. In contrast to the declining economic performance of CE and the consistently sub-optimal economic yield for TAC, the POMDP strategy continues to generate economic yields at approximately the optimal level (as attained by CE in the absence of measurement error) despite the increasingly uncertain measurements. This demonstrates that the reduced risk of stock collapse and higher average stock biomass attained by the POMDP strategy in the presence of high measurement error, Figure 1) is not the result of a trivial reduction in harvesting under all circumstances, but rather, evidence of a more nuanced strategy that manages to account for the uncertainty in measurement while maintaining a reasonable harvest.

The POMDP strategy is also robust to overestimation of the level of measurement uncertainty. In the simulation results shown in Figures 1 and 2, we have assumed that the level of measurement uncertainty σ_m was known, and saw that ignoring this uncertainty (as the CE policy does) has significant negative impact on ecological and economic outcomes. If the level of measurement uncertainty is not known precisely, *overestimating* the measurement error while using the POMDP strategy provides a precautionary approach that can nevertheless achieve nearly optimal ecological and economic outcomes. Figure 3 summarizes the

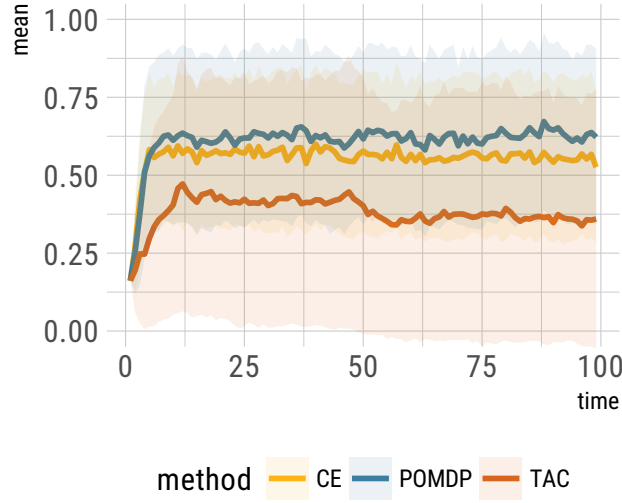


Figure 3: Average fish biomass when the POMDP strategy assumes a high level of measurement uncertainty, while simulations reflect perfect measurements. Despite this overestimation, biomass under POMDP strategy closely tracks the optimal biomass under CE in the absence of measurement error.

results of the same simulations as before, but under the scenario in which the POMDP approach assumes a measurement error of $\sigma_m = 0.15$, when in fact all simulated measurements are made without error ($\sigma_m = 0$). This represents an extreme case of overestimating the measurement error. Stochastic growth remains the same as before, $\sigma_g = 0.15$, and simulations under TAC and CE policies are shown for comparison. When measurement error is absent, Reed’s proofs hold and the CE strategy is optimal. Figure 3 shows that the POMDP outcomes track almost exactly the CE outcomes despite the misplaced assumption that measurements are quite poor. (As we have already seen, the TAC strategy is insufficiently cautious for this level of stochasticity in recruitment, resulting in over-exploitation and long-term decline). The ability of the POMDP solution to perform nearly optimally even when significantly overestimating the level of measurement uncertainty contrasts sharply to the significant declines from ignoring measurement uncertainty seen in the CE solutions in Figure 1. This demonstrates that while we may not know precisely the level of measurement error, we achieve far better outcomes overestimating measurement uncertainty than underestimating it. This also underscores the observation that POMDP policy is quite robust to the details of the uncertainty. We can get a better understanding for the performance of POMDP in these simulations by looking more closely at how any individual decision under a POMDP strategy compares to the action chosen by the current alternatives.

To better understand the differences in performance of these strategies and resolve the paradox of uncertainty, we must take a closer look at how the specific action recommended by each policy compares given the same observation. Figure 4 shows the action taken by each strategy in response to a measurement




Figure 4: Comparison of the harvest and escapement policies under each strategy as a function of the observed stock size. Escapement refers to the expected fraction of fish left in the sea, $x - h$, while harvest refers to the target catch; two different conventions for plotting the same action. Uniquely, the POMDP policy plot will depend not only on the observed stock size, but also depends on prior information (bottom panel) as determined from any prior observations and actions. Depending on this prior it may harvest less or more than the other policies given an identical observation.

300 of the stock size (biomass estimate). These plots show, for any possible observations, which strategy will
 attempt to harvest most and which will attempt to harvest least. For comparison purposes, we plot policies
 both in terms of expected escapement, $S = Y - H$ as is typical in the optimal control literature, (e.g. Reed
 303 1979; Clark and Kirkwood 1986; Sethi et al. 2005), and also directly in terms of the harvest quota H . These
 plots clearly illustrate the contrast between the constant escapement (CE) strategy and the precautionary
 Total Allowable Catch (TAC) strategy: CE sets harvest strictly to zero for stocks estimated at biomass below
 306 B_{MSY} , while TAC permits a modest harvest of even very small stocks. (Comparable plots with MSY can be
 found in the Supplemental Material). In principle, the CE policy could depend on σ_g , but as Reed (1979)
 proved, the constant escapement level with stochasticity is the same as in the deterministic case, $S = D$,
 309 unless the noise level is quite high (comparison plots in Supplemental material.) In contrast to this, our
 POMDP plots depend very much on both the choice of σ_m and σ_g . If $\sigma_m = 0$, they reduce exactly to the CE
 solution. Here we show the corresponding policies for POMDP solutions focusing on $\sigma_g = 0.15$ as above, with
 312 a modest measurement uncertainty $\sigma_m = 0.1$. Alternate combinations of measurement uncertainty can be
 found in the appendix but do not change the general pattern. In addition to an explicit dependence on the
 measurement error, our POMDP solutions depend on another piece of information: any prior observations of
 315 the stock size.

Differences in attention to prior beliefs, determined by prior observations, drive the differences between
 the POMDP strategy and the other strategies and can resolve the uncertainty paradox. While the catch
 318 quota under both TAC and CE strategies can be completely determined given the most recent observation of
 the stock size by using the policy curves shown in Figure 4, this is not the case for the POMDP approach.
 This fundamental difference is key to understanding the difference in performance and resolving the paradox
 321 of uncertainty. The POMDP policy cannot be specified by the most recent observation alone. Instead, the
 POMDP policy depends on all prior observations, not just the most recent. The reason for this complexity
 comes from the Markov property. Observations of the state in the perfectly observed system satisfy the
 324 Markov property: once we have measured the current biomass exactly, we cannot get any better estimate of
 the current stock size by studying older measurements. When measurements are uncertain this is no longer
 the case: intuitively, by comparing the most recent measurement to previous observations we may be able to
 327 infer when any given measurement is unusually high or unusually low. POMDP formalizes this intuition by
 capturing the information from all previous observations into a *prior belief*. This prior belief is updated after
 every subsequent action and observation in accordance with Bayes Law. The mechanics of this process are
 330 well documented in the extensive literature on POMDPs (e.g. Smallwood and Sondik 1973; Sondik 1978;

Kurniawati et al. 2008; Williams 2011), but for our purposes it is sufficient to observe how this evolving prior belief serves to continually adjust the POMDP policy. Supplemental Figure S9 (Appendix A) illustrates how the prior belief evolves in response to subsequent observations and actions over the course of an individual simulation.

Figure 4 shows two separate policy curves (in terms of harvest and escapement) for the same POMDP solution given two different prior belief distributions (panel 3, priors). While both priors shown express considerable uncertainty about the precise stock size prior to the most recent observation, the lower prior is centered at a value $\frac{1}{4}K$ stock size, while the high prior is centered at a value of $\frac{3}{4}K$, relative to a (post-harvest, before observation) target size of $B_{MSY} = K/2$. The POMDP policy is determined by the combination of this prior information and the most recent observation, as indicated by the two different POMDP curves for harvest/escapement shown corresponding to the different priors. As with the other policies, the higher the most recent observation (x axis) the higher the POMDP recommended harvest. Yet unlike the alternative strategies, the POMDP solution always reflects the prior information. Consequently, relative to constant escapement (that is, no measurement error), the POMDP with low prior starts harvesting only at higher stock sizes and always harvests less. In contrast under the high prior, the POMDP always harvests at the same or higher level than the constant escapement solution.

A resolution to the paradox.

Herein lies our resolution to the paradox of uncertainty. Previous work created this paradox by suggesting that increased harvest rates (decreased target escapement) would often be the rational response to increased uncertainty. The exact solutions from the POMDP reveal that this is only an accident of the assumptions: it is indeed true that under certain circumstances, harvest levels should increase relative to the case of no uncertainty, but *only when prior knowledge suggests the stock size should be much higher than the most recent estimate would suggest*. In the POMDP solution, all information must be put into its historical (and constantly updated) context. When a measurement roughly matches the expectation of this prior context, Figure 4 shows that the optimal response from POMDP is roughly comparable or slightly more cautious to the harvest under no uncertainty, and not more aggressive as the paradox would suggest. Measurements that exceed expectations are tempered with some skepticism: while the CE solution is willing to meet a high stock measurement with a large harvest, the POMDP solutions increase harvest more cautiously.

This difference between underestimating and overestimating accounts for the poor performance of the CE solution under large measurement uncertainty, where it over-harvests whenever measurements are too large. Even though underestimating is equally likely under the measurement uncertainty model, sooner or




Figure 5: Comparisons of harvest level that would be recommended by policies considered here relative to historical harvest levels in two commercial fish stocks. POMDP solution assumes a measurement error of 10%.

later a run of “heads”, a sequence of overestimations relative to the true stock, can drive stocks to very low levels where the chance of local extinction becomes possible. It is precisely this asymmetry: that too much over-harvesting leads to an irreversible state of extinction, while too much under-harvesting is always reversible (modulo some lost revenue) that lay behind Clark’s original intuition that there was something fishy about Reed’s result that $S = D$: that uncertainty required no extra caution. Constant escapement is particularly susceptible to over-harvesting starting from stock sizes much higher than B_{MSY} since its bang-bang optimal solution attempts to bring stocks back to target level as fast as possible. CE is not so vulnerable to collapse from small stock sizes, since it shuts down all harvests once estimates fall below B_{MSY} . Stock collapse under TAC can also be driven by such a string of heads, but is unlikely at high stock sizes since harvest never exceeds H_{MSY} . In both cases, some measurement uncertainty interacts with inherent stochasticity, which provides a continued source of variation to sizes above B_{MSY} , where CE strategy is most vulnerable, and below B_{MSY} , where TAC is most vulnerable.

Historical examples and implementation.

So far we have focused on simulation and an examination of the relative policies under different prior

information. Figure 5 compares the harvest level that would have been recommended by each of the strategies we have considered here against the historically observed catch recorded for two commercially fished stocks: Argentine Hake and European Plaice. Historical estimates of biomass and catch are taken from the R.A. Myers Legacy Stock Assessment database (Ricard et al. 2011). Posterior distributions for parameters for the Gordon-Schaefer model are estimated with uncertainty through Markov Chain Monte Carlo using `nimble` (de Valpine et al. 2017) on the historical data, illustrating how this process might be done in more complex models as well (details and code in Supplement, Appendix B). B_{MSY} and corresponding TAC and CE policies are calculated based on posterior mean estimates of the model parameters, along with POMDP solution assuming a 10% measurement error rate. Though measurement error could be estimated directly from the raw data, this would not reflect the true measurement uncertainty arising from the stock assessment process. Each strategy is then compared to historical observations to determine a recommended harvest. As we have seen, the TAC and CE harvest policies are uniquely determined by the observed stock size relative to B_{MSY} , but the POMDP policy must be re-calculated each time step to reflect both the prior observation and prior action.

These historical examples provide another useful lens to compare how the different strategies respond in the face of fluctuations in real data, rather than model simulations. In both stocks, historical catch almost always exceeds that recommended by any of our strategies, falling closer to the MSY value estimated (see Appendix B). In Argentine Hake, the persistent declines result in both POMDP and CE strategies quickly closing down the fishery in an attempt to let the stock recover to a more productive target biomass, while TAC persists with merely a reduced harvest. A small recovery of biomass in 1995 is met by an immediate uptick in harvest under CE, while the POMDP response is more conservative. The example of European Plaice illustrates more volatility in the stock size, revealing further differences in the strategies. Despite this volatility, the TAC level remains relatively level across all five decades, rising or dipping only slightly with changes in stock size. In contrast, the CE harvest tracks this volatility almost exactly. Here, the POMDP solution falls in between these extremes, almost mirroring which-ever of the two policies is more conservative: the POMDP solution matches the dips in harvest taken by CE to allow the stock to recover most quickly, but does not track the doubling of harvests recommended by CE in the mid 1980, increasingly only to the more modest harvest recommended by the TAC policy. Once again, we see that POMDP solution provides a more consistently precautionary policy than either the over-simplified optimal control solution of CE or the more rule-of-thumb approach represented by TAC. Though we have seen that the POMDP solution need not always be more conservative (lower harvest) than these strategies, it makes use of the all prior observations

to tune the level of caution appropriately.

Discussion

Using modern algorithms we have been able to crack the nut of measurement uncertainty in the optimal management of marine fisheries and debunk long-standing paradoxes that increased measurement uncertainty should be met with larger harvests. We have demonstrated that both existing optimal control approaches and existing precautionary approaches result in over-exploitation and long term declines of fish stocks, as well as reduced economic value, while in contrast, the POMDP approach is able to deliver both ecological and economic recovery despite uncertainty in measurements. These results have important implications for both global fisheries and ecological management more generally. In marine fisheries, evidence of such over-exploitation has already been well documented (Worm et al. 2006, 2009; Costello et al. 2016). Our results suggest that measurement uncertainty may contribute to this pattern, as such errors lead to over-exploitation under current best practices such as MSY, TAC or CE, which have previously been thought to be sufficient to ensure the recovery of stocks world-wide (Costello et al. 2016).

Our results also have important consequences for ecological management more generally. The study and application of management strategies such as MSY and CE underpin much of conservation and natural resource management in forest management and many other terrestrial ecosystems (e.g. Clark and Mangel 2000). Decision making under uncertainty is a central challenge for ecology and conservation biology (Polasky et al. 2011), which we have long attempted to address through either simplified optimal control models or more heuristic approaches such as the precautionary principle and resilience thinking (Fischer et al. 2009). We have illustrated how both of these approaches can fail to provide an adequate strategy when faced with measurement uncertainty. To avoid this pitfall, a strategy must be dynamic to make best use of the available information without assuming that information is perfect. The simplistic optimal control approach of CE fails by placing too much confidence in measurements even when they are significantly at odds with prior expectations. Similarly, the TAC approach also fails due to its reliance on a simple one-size-fits-all strategy, not reducing catch quotas sufficiently when stocks are too low but also failing to increase them when stocks are higher. It is precisely this ability to adapt the policy in response to available information, as illustrated in Figure 4, that makes the POMDP approach successful. Nevertheless, it may be tempting to see these results as either trivial or impractical, but a closer inspection reveals they are neither.

First, it may seem obvious that POMDP would out-perform more primitive optimal control solutions such as CE, but this was hardly a forgone conclusion. Indeed, the previous example of Reed (1979) had

proven that the complex stochastic dynamic programming solutions required to accommodate stochastic
 438 dynamics would perform no better than the simple maximization based purely on the deterministic model. If
 stochasticity doesn't matter, it is easy to believe that measurement error would be minor at best. In fact, we
 find that once we account for measurement error, the effects of both stochasticity and measurement have
 441 a significant influence over the optimal policy. It is similarly tempting to dismiss the better performance
 of POMDP against a more rule-of-thumb^[1] based strategy like the TAC scenario considered here (fishing
 at 80% MSY), but once again this was not a forgone conclusion. While a better economic return might
 444 reasonably be expected of the optimization approach, it was not obvious that POMDP would also be at
 times more cautious than the rule of thumb, able to deliver both higher economic returns by harvesting more
 when it was safe to do so, and higher ecological recovery by harvesting even less than TAC when necessary.
 447 These results highlight the difficulty of finding effective heuristic rules and underscore the power and utility
 of more sophisticated strategies such as POMDP which can make the best use of all available information.

¹: Given the sophisticated policies often in place to determine actual TAC levels, one might object to
 450 our characterization of this as merely a rule-of-thumb. We use this term merely to distinguish it from the
 optimal control approach – that despite technical precision, TAC adjustments are set by a heuristic instead
 of the result of some dynamic optimization algorithm as in CE and POMDP.

Second, we must ask if the policies derived from POMDP are practical. As with constant escapement,
 453 the POMDP solutions we considered here result in more year-to-year variation in harvest, and include the
 possibility of fishing moratorium to allow stocks to recover – decisions that many view as economically or
 456 socially unacceptable. While resolving this debate is beyond the scope of this paper, a few observations are
 relevant: (A) Stocks already managed under constant escapement, such as US salmon fisheries, are already
 subject to the possibility of these moratoriums. (B) It is straight-forward to introduce constraints on the
 459 minimal acceptable harvest level and/or include the additional costs of closing a fishery in the utility function
 optimized by POMDP. (C) Fisheries today are not managed at either their maximum ecological or economic
 potential. Taken as a whole, most ocean fisheries yield a net negative economic return (World Bank and FAO
 462 2009), with fishing persisting only with the support of government subsidies (Arnason 2012), meanwhile some
 68% of global fisheries are in poor biological condition (Costello et al. 2016). While the causes for this are
 both complex and numerous, we have seen that both the economic inefficiency and ecological health suffer
 465 under TAC/MSY and CE based management. Arnason (2012) and Costello et al. (2016) have both argued
 that a rights-based approach to fisheries management (RBFM) would lead to self adoption of precisely such
 optimal policies.

Third: are the examples considered here too simple? The strategies we have compared against underpin both current practice and management aspirations. Together, MSY and TAC-based management reflects much of modern best practice as it is understood in fisheries management and elsewhere. While real-world application of these approaches frequently involves more complex models, the models we have considered here are the very ones under which these strategies were designed and tested, and there is little reason to believe they would perform better under more complex cases such as age-structured models that do not always meet those assumptions (e.g. Holden and Conrad 2015). The relatively poor ecological and economic performance of these common best-practices relative to the alternative approach of POMDP is thus reason enough for concern. This issue is further underscored by the reality of both long-running ecological deterioration (Worm et al. 2006, 2009) and economic costs of global fisheries.

Fourth: is the approach taken here too complex to be feasible? The importance of simple and effective rules of thumb in conservation management has been well documented (e.g. Chadès et al. 2011). Sophisticated and computationally intensive approaches such as POMDP may seem improbable at a time when many areas of natural resource management, even simple, rule-of-thumb based methods struggle to take root. And yet, in this same age we rely on far more complex and computationally intensive algorithms to perform far more trivial tasks such as determining what advertisements we see. Surely a planet that can afford such complex optimization for how advertising space is allocated on a small screen can one day manage its natural resources and conservation efforts with such tools?

Future directions

Several limitations that have been studied in fully observed (MDP) optimal decision problems, such as parameter uncertainty (e.g. Ludwig and Walters 1982), model uncertainty (Williams 2001; Boettiger et al. 2015) and adaptive management (e.g. Walters and Hilborn 1976) remain largely open challenges for partially observed systems. Future work could extend this analysis to more complex models, such as those with age structure, as (Holden and Conrad 2015) does for the fully observed case. Another limiting assumption common to MDPs and POMDPs is that of stationary dynamics: that the population dynamics equation itself is not changing over time. In reality, forces such as climate change and other forms of environmental variations violate this assumption (Britten et al. 2017). Direct approaches such as Fackler and Pacifici (2014)’s adaptation of Mixed Observability MDP (Ong et al. 2010) do not scale to the number of states and actions considered here. A value of information (VOI) analysis for POMDP (e.g. Johnson and Williams 2015; Memarzadeh and Pozzi 2016), could identify when it is worthwhile to actively reduce measurement error.

References

Arnason, R. 2012. Property rights in fisheries: How much can individual transferable quotas accomplish? *Review of Environmental Economics and Policy* 6:217–236.

Beverton, R., and S. Holt. 1957. *On the Dynamics of Exploited Fish Populations*. Chapman; Hall, London.

Boettiger, C., M. Mangel, and S. Munch. 2015. Avoiding tipping points in fisheries management through Gaussian process dynamic programming. *Proceedings of the Royal Society B: Biological Sciences* 282:20141631–20141631.

Boettiger, C., J. Ooms, and M. Memarzadeh. 2018. Sarsop: Approximate pomdp planning software.

Botsford, L. W., J. C. Castilla, and C. Peterson. 1997. The management of fisheries and marine ecosystems. *Science* 277:509–515.

Britten, G., M. Dowd, L. Canary, and B. Worm. 2017. Extended fisheries recovery timelines in a changing environment. *Nature Communications*.

Chadès, I., T. G. Martin, S. Nicol, M. A. Burgman, H. P. Possingham, and Y. M. Buckley. 2011. General rules for managing and surveying networks of pests, diseases, and endangered species. *Proceedings of the National Academy of Sciences* 108:8323–8.

Chadès, I., E. McDonald-Madden, M. a McCarthy, B. Wintle, M. Linkie, and H. P. Possingham. 2008. When to stop managing or surveying cryptic threatened species. *Proceedings of the National Academy of Sciences* 105:13936–40.

Clark, C. W. 1973. Profit maximization and the extinction of animal species. *Journal of Political Economy* 81:950–961.

Clark, C. W. 1990. *Mathematical Bioeconomics: The Optimal Management of Renewable Resources*, 2nd Edition. Wiley-Interscience.

Clark, C. W., and G. P. Kirkwood. 1986. On uncertain renewable resource stocks: Optimal harvest policies and the value of stock surveys. *Journal of Environmental Economics and Management* 13:235–244.

Clark, C. W., and M. Mangel. 2000. *Dynamic state variable models in ecology*. Oxford University Press, Oxford.

Costello, C., D. Ovando, T. Clavelle, C. K. Strauss, R. Hilborn, M. C. Melnychuk, T. A. Branch, et al. 2016. Global fishery prospects under contrasting management regimes. *Proceedings of the National Academy of Sciences* 113:5125–5129.

de Valpine, P., D. Turek, C. J. Paciorek, C. Anderson-Bergman, T. Duncan, and R. Bodik. 2017.

Programming With Models: Writing Statistical Algorithms for General Model Structures With NIMBLE.
Journal of Computational and Graphical Statistics 26:403–413.

531 Dichmont, C. M., R. A. Deng, A. E. Punt, J. Brodziak, Y.-J. Chang, J. M. Cope, J. N. Ianelli, et al.
2016. A review of stock assessment packages in the united states. Fisheries Research 183:447–460.

Engen, S., R. Lande, and B. Sæther. 1997. Harvesting strategies for fluctuating populations based on
534 uncertain population estimates. Journal of Theoretical Biology 186:201–212.

Fackler, P., and R. Haight. 2014. Monitoring as a partially observable decision problem. Resource and
Energy Economics 37:226–241.

537 Fackler, P., and K. Pacifici. 2014. Addressing structural and observational uncertainty in resource
management. Environmental Management 133:27–36.

Fischer, J., G. D. Peterson, T. a Gardner, L. J. Gordon, I. Fazey, T. Elmqvist, A. Felton, et al. 2009.
540 Integrating resilience thinking and optimisation for conservation. Trends in ecology & evolution 24:549–54.

Gordon, H. S., and C. Press. 1954. The Economic Theory of a Common-Property Resource: The Fishery.
Journal of Political Economy 62:124–142.

543 Hilborn, R. 2010. Pretty Good Yield and exploited fishes. Marine Policy 34:193–196.

Holden, M., and J. Conrad. 2015. Optimal escapement in stage-structured fisheries with environmental
stochasticity. Mathematical biosciences 269:76–85.

546 Holling, C. S. 1973. Resilience and Stability of Ecological Systems. Annual Review of Ecology and
Systematics 4:1–23.

Johnson, F. A., and B. K. Williams. 2015. A Decision-Analytic Approach to Adaptive Resource
549 Management. Pages 61–84 in C. R. Allen and A. S. Garmestani, eds. Adaptive management of social-
ecological systems. Springer Netherlands, Dordrecht.

Kaelbling, L. P., M. L. Littman, and A. R. Cassandra. 1998. Planning and Acting in Partially Observable
552 Stochastic Domains. Artificial Intelligence 101:99–134.

Kurniawati, H., D. Hsu, and W. S. Lee. 2008. SARSOP : Efficient Point-Based POMDP Planning by
Approximating Optimally Reachable Belief Spaces. Proceedings of Robotics: Science and Systems IV.

555 Larkin, P. a. 1977. An Epitaph for the Concept of Maximum Sustained Yield. Transactions of the
American Fisheries Society 106:1–11.

Levin, S. A., and J. Lubchenco. 2008. Resilience, Robustness, and Marine Ecosystem-based Management.
558 BioScience 58:27–32.

Ludwig, D., and C. Walters. 1981. Measurement errors and uncertainty in parameter estimates for stock

and recruitment. *Journal of Canadian Fisheries and Aquatic Sciences* 38:711–720.

Ludwig, D., and C. J. Walters. 1982. Optimal harvesting with imprecise parameter estimates. *Ecological Modelling* 14:273–292.

Mace, P. M. 2001. A new role for MSY in single-species and ecosystem approaches to fisheries stock assessment and management. *Fish and Fisheries* 2:2–32.

Mangel, M., and C. W. Clark. 1988. *Dynamic Modeling in Behavioral Ecology*. (J. Krebs & T. Clutton-Brock, eds.). Princeton University Press, Princeton.

Marescot, L., G. Chapron, I. Chadès, P. L. Fackler, C. Duchamp, E. Marboutin, and O. Gimenez. 2013. Complex decisions made simple: a primer on stochastic dynamic programming. *Methods in Ecology and Evolution* n/a–n/a.

May, R. M. 1977. Thresholds and breakpoints in ecosystems with a multiplicity of stable states. *Nature* 269:471–477.

Memarzadeh, M., and M. Pozzi. 2016. Value of information in sequential decision making: component inspection, permanent monitoring and system-level scheduling. *Reliability Engineering & System Safety* 154:137–151.

Moxnes, E. 2003. Uncertain measurements of renewable resources: approximations, harvesting policies and value of accuracy. *Journal of Environmental Economics and Management* 45:85–108.

Ong, S., S. Png, D. Hsu, and W. Lee. 2010. Planning under uncertainty for robotic tasks with mixed observability. *The International Journal of Robotics Research* 29:1053–1068.

Petchey, O. L., M. Pontarp, T. M. Massie, S. Kéfi, A. Ozgul, M. Weilenmann, G. M. Palamara, et al. 2015. The ecological forecast horizon, and examples of its uses and determinants. *Ecology Letters* 18:597–611.

Pineau, J., G. Gordon, and S. Thrun. 2003. Point-based value iteration: An anytime algorithm for POMDPs. *IJCAI International Joint Conference on Artificial Intelligence* 1025–1030.

Polasky, S., S. R. Carpenter, C. Folke, and B. Keeler. 2011. Decision-making under great uncertainty: environmental management in an era of global change. *Trends in ecology & evolution* 1–7.

Reed, W. J. 1979. Optimal escapement levels in stochastic and deterministic harvesting models. *Journal of Environmental Economics and Management* 6:350–363.

Ricard, D., C. Minto, O. Jensen, and J. Baum. 2011. Examining the knowledge base and status of commercially exploited marine species with the RAM Legacy Stock Assessment Database. *Fish and Fisheries* 13:380–398.

Roughgarden, J. E., and F. Smith. 1996. Why fisheries collapse and what to do about it. *Proceedings of*

- 591 the National Academy of Sciences 93:5078.
- Schaefer, M. B. 1954. Some aspects of the dynamics of populations important to the management of the commercial marine fisheries. *Bulletin of the Inter-American Tropical Tuna Commission* 1:27–56.
- 594 Sethi, G., C. Costello, A. Fisher, M. Hanemann, and L. Karp. 2005. Fishery management under multiple uncertainty. *Journal of Environmental Economics and Management* 50:300–318.
- Smallwood, R. D., and E. J. Sondik. 1973. The Optimal Control of Partially Observable Markov Decision
597 Processes over a Finite Horizon.
- Sondik, E. J. 1978. The Optimal Control of Partially Observable Markov Processes Over the Infinite Horizon : Discounted Costs. *Operations Research* 26:282–304.
- 600 Walters, C. J., and R. Hilborn. 1976. Adaptive Control of Fishing Systems. *Journal of the Fisheries Research Board of Canada* 33:145–159.
- . 1978. Ecological Optimization and Adaptive Management. *Annual Review of Ecology and*
603 *Systematics* 9:157–188.
- Williams, B. K. 2001. Uncertainty , learning , and the optimal management of wildlife. *Environmental and Ecological Statistics* 8:269–288.
- 606 ———. 2011. Resolving structural uncertainty in natural resources management using POMDP approaches. *Ecological Modelling* 222:1092–1102.
- World Bank and FAO. 2009. The Sunken Billions. The Economic Justification for Fisheries Reform.
- 609 Worm, B., E. B. Barbier, N. Beaumont, J. E. Duffy, C. Folke, B. S. Halpern, J. B. C. Jackson, et al. 2006. Impacts of biodiversity loss on ocean ecosystem services. *Science (New York, N.Y.)* 314:787–90.
- Worm, B., R. Hilborn, J. K. Baum, T. A. Branch, J. S. Collie, C. Costello, M. J. Fogarty, et al. 2009.
612 Rebuilding global fisheries. *Science (New York, N.Y.)* 325:578–85.



[Click here to access/download](#)

Other (Video, Excel, large data files)
appendixA.pdf



[Click here to access/download](#)

Other (Video, Excel, large data files)
appendixB.pdf

